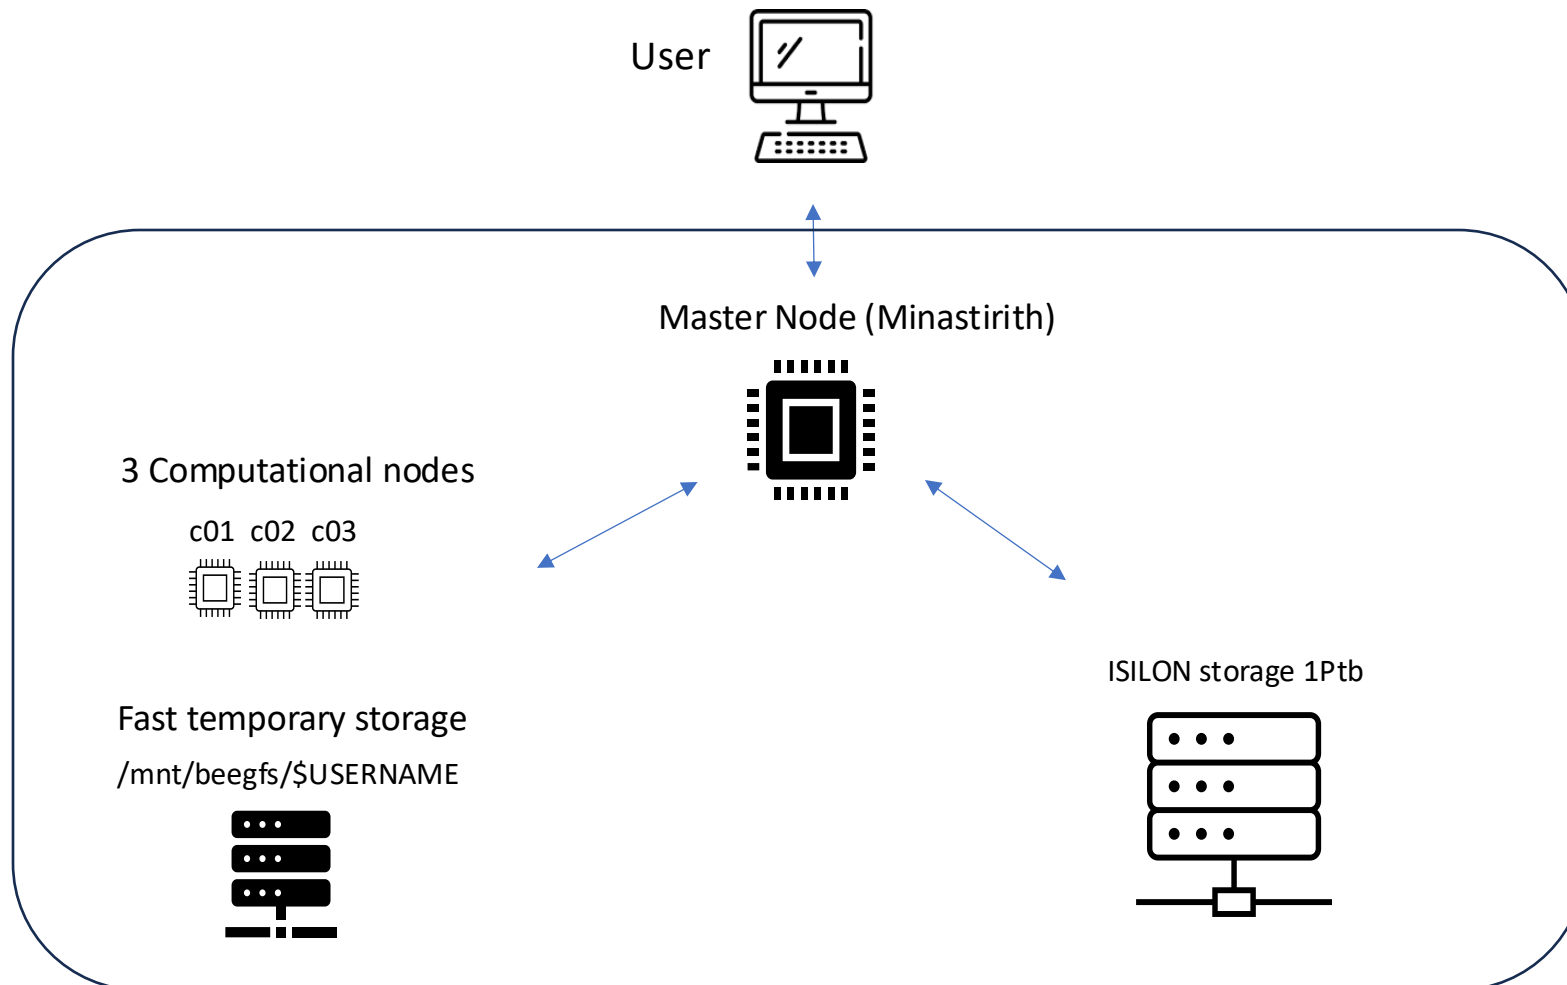


# Computations with the IJC High-performance computing cluster



Angelika Merkel (Head of Bioinformatics Core Facility)  
09/04/2025

# IJC HPC cluster



# IJC cluster components




- Master Node (Minastirith)

12 cores 128Gb Ram

= connects user, computational nodes (+ flash storage), ISILON

USE: to connect, submit slurm jobs

-> Shared resource, do not run your jobs here!



- Computational nodes

c01 128 cores 1Tb Ram


c02 128 cores 1Tb Ram

c03 128 cores 1Tb Ram GPU

= computing (job execution via SLURM)

USE: to execute parallel or high-mem jobs

-> Shared resource, use wisely!



- Flash storage '/mnt/beegfs/'

100Tb SSD memory

= flash memory attached to computational nodes for fast read/write access

USE: to temporarily store data

-> Shared resource, use wisely!



# Connecting to the cluster



You can connect to the cluster (master node = minastirith) in several ways:

1. From anywhere with a browser via VPN Portal: <https://vpn.carrerasresearch.org> (bookmark 'Minastirith')
2. From your machine (Linux/Mac) with a terminal:  
`$ ssh username@minastirith`
3. From your machine using a ssh client (e.g. [PUTTy for windows](#) ) connect directly to the host:  
`minastirith`

To connect from outside the IJC network you need to [connect to the SSHlogin node](#) first.

# Enter minastirth



\$ ssh amerkel@minastirth

```
[amerkel@IJC20249 ~]$ ssh minastirth
Last login: Mon Jun 10 17:06:04 2024 from 84.88.90.193

      o8o      -      o8o      o8o      .      oooo
      .o8      .o8      .o8      .o8      888
ooo. .oo. .oo. oooo ooo. .oo. .oooo. .oooo.o 888 oooo oood8b oooo 888 888
'888P'Y88bP'Y88b '888 '888P'Y88b `P )88b d88( "8 .o888oo '888 oooo d8b '888 .o888oo 888 .oo.
888 888 888 888 888 888 .oP'888 ``Y88b. 888 888 888 888 888 888 888 888
888 888 888 888 888 888 d8( 888 o. )88b 888 . 888 888 888 888 . 888 888
o888o o888o o888o o888o o888o o888o Y888 -8o 8 888P 888 o888o d888b o888o 888 o888o o888o

Master AMD EPYC 7272 12 cores 128Gb Ram *****
Node01 AMD EPYC 7702 128 cores 1Tb Ram *****
Node02 AMD EPYC 7702 128 cores 1Tb Ram *****
Node03 AMD EPYC 7702 128 cores 1Tb Ram Nvidia A40 GPU

Please remember:

1.- Save your stuff before it gets deleted! <3

    BEEGFS data (/mnt/beeufs) will be periodically removed.
    Transfer your data to /ijc/LABS or /ijc/USERS !

    Example: $ rsync -rlv --checksum source_folder destination_folder

2.- Do not launch jobs on the master node, use sbatch files or salloc. Interactive jobs on minastirth will be killed!

3.- Check performance of your finished jobs!

    Example: $ seff [jobid]

Thank you all!

Partition Name  Max Time  Default Time  Max CPU x Node  Max Mem x CPU  Def Mem x CPU  Max Mem x Node
=====
fast            2 hours   1h            64             4000M          2000M          256000M
highMem         1 Day     1h            8              -              -              800000M
hpc             3 Days    4h            80             200000M        4000M          -

Max run Cpu Per User: 48

show personal statistics
=====
module load reportseff
reportseff --since d=2 --format "+user,jobname,start,ncpus,reqmem" -u amerkel
[amerkel@minastirth ~]$
```



# Working with the HPC

input → software → instructions → resources → output

# Using the software



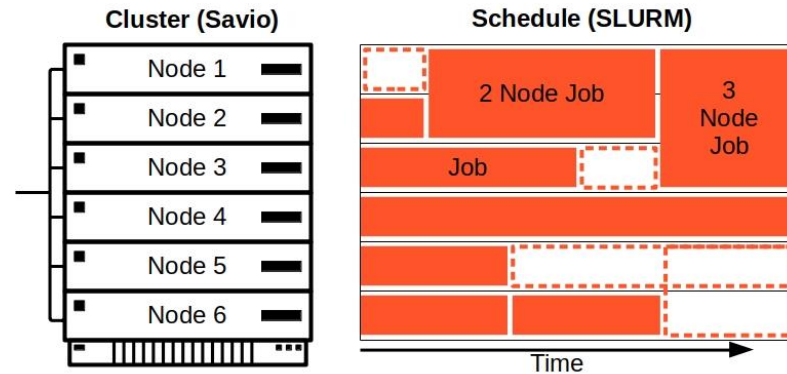
Available as modules

| Description               | Command                         |
|---------------------------|---------------------------------|
| see all available modules | <code>module avail</code>       |
| load/unload a module      | <code>module load/unload</code> |
| see all loaded modules    | <code>module list</code>        |
| unload all modules        | <code>module purge</code>       |

# Executing a job



All tasks (jobs) executed on the cluster computational nodes are managed by SLURM (the scheduler)



<https://docs-research-it.berkeley.edu/services/high-performance-computing/user-guide/>

SLURM schedules each job based on available resources (CPU, memory, nodes, execution time, etc)



# SLURM script



`./my_clusterjob.sh # basic serial job`

```
#!/bin/bash

# SLURM arguments
#SBATCH --job-name=job_serial      # Job name
#SBATCH --cpus-per-task=1          # Run on a single CPU
#SBATCH --mem=4gb                  # Job memory request
#SBATCH --time=00:10:00            # Time limit hrs:min:sec
#SBATCH --output=job_%j.log        # Standard output and error log

# load software
module load R

# message something
echo "Running R script on a single CPU core"

# run
Rscript myscript.r
```

`%j = jobid`

# Slurm Commands



| Command                              | Description   |
|--------------------------------------|---|
| <code>sbatch my_clusterjob.sh</code> | Submit job for execution  |
| <code>squeue</code>                  | Show the actual job queue   |
| <code>scancel jobid</code>           | Cancel a job  |
| <code>sacct -j jobid</code>          | Job accounting infos  |
| <code>seff jobid</code>              | Report on the efficiency of a job's cpu and memory utilization (after the job has finished) |
| <code>salloc</code>                  | Allocate resource for an interactive shell  |

```
module load reportseff
reportseff --since d=2 --format "+user,jobname,start,ncpus,reqmem" -u username
```

# Deploy interactive shell



You can deploy an interactive shell (with specified resource) on a computational node for testing:


```
$ salloc -c 4
```

```
[amerkel@minastirith ~]$ salloc
salloc: Granted job allocation 3216713
salloc: Waiting for resource configuration
salloc: Nodes c02 are ready for job
[amerkel@c02 ~]$
```


# Parallelization increases speed



1. Single input + single task



2. Multiple input + Single task



# Job arrays



```
#!/bin/bash

# SLURM arguments
#SBATCH --job-name=job_array      # Job name
#SBATCH --ntasks=1               # Run a single task
#SBATCH --mem=4gb                 # Job Memory
#SBATCH --time=00:05:00          # Time limit hrs:min:sec
#SBATCH --output=array_%A-%a.log # Standard output and error log
#SBATCH --array=1-5              # Array range

# message something
echo "This is task" $SLURM_ARRAY_TASK_ID

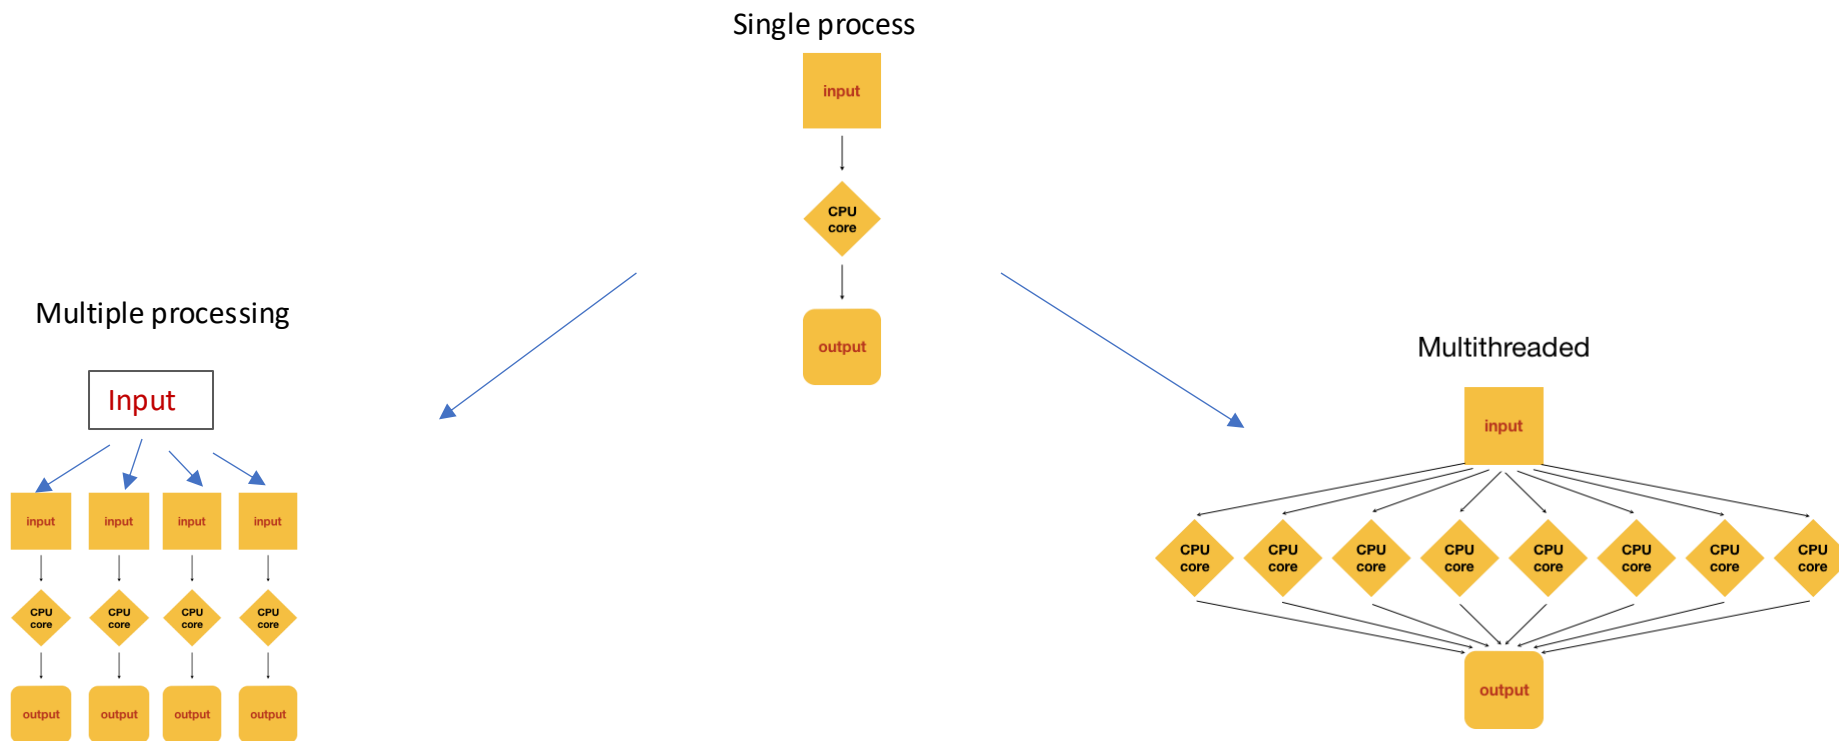
# run
infile=$(ls *.txt | sed -n ${SLURM_ARRAY_TASK_ID}p) # multiple input files (*.txt)

./myscript.sh $infile                                # run myscript for each input file
```

Job arrays allow to manage and submit similar jobs or execute the same job over multiple inputs (samples/data sets/parameter sets)

- Multiple jobs are created/submitted from one script, each with a different `$SLURM_ARRAY_TASK_ID` (=build-in variable)
- `$SLURM_ARRAY_TASK_ID` takes value specified with the array range


# Parallelization increases speed



# Multi-processing



Multiple processes




In **multi-processing** multiple processes are executed on multiple CPUs:

- `#SBATCH --ntask=number_of_processes`

# Multi-threading



Multithreaded



In **multi-threading** a single process is executed with multiple threads (e.g. multiple code segments):

1. Shared-memory multi-threading application (threaded, OpenMP, PTHREADS) can use multiple cpus but only on a single node:

- #SBATCH --nodes=1
- #SBATCH --ntask=1
- #SBATCH --cpus-per-task=number\_of\_threads

2. OpenMPI applications allow to share threads across nodes:

- #SBATCH --ntask=1
- #SBATCH --cpus-per-task=number\_of\_threads



# SLURM partitions (queues)



There are multiple partitions on the cluster depending to which jobs can be submitted depending on the required resources

#SBATCH --partition

```
Partition Name  Max Time  Default Time  Max CPU x Node  Max Mem x CPU  Def Mem x CPU  Max Mem x Node
=====
fast           2 hours   1h            64              4000M          2000M          256000M
highMem        1 Day     1h            8               -              -              800000M
hpc            3 Days    4h            80              200000M        3000M          -

Max Cpu by job: 16
```

**\*NOTE:** If the resource requirement is not specified, default value = max value

# Do's & Don't's

## DO

- only request resources that you need (more CPUs or memories don't make your job faster!)
- check if the resources you requested were efficiently used (use `seff -j jobid` or similar)
- check if your job should be submitted to a special queue
- submit resource intensive jobs at low peak times (over night or the weekend)
- break large jobs down into smaller ones
- optimize your code

## DON'T

- ask for unnecessary resources (+ 10/15% time or memory are usually sufficient)
- Do not execute any resource intensive task on the master node, use an interactive shell instead
- forget to close your interactive session

# How to get help

- contact **IT** for help with resource requirements, software and permissions
- contact **BIT** for help with coding and workflows
- check the **documentation** @:  
<https://howto.carrerasresearch.org/>  
<https://vpn.carrerasresearch.org/> > wordpress
- check the BIT workshop pages @:  
[https://ijcbit.github.io/Workshops/Trainings/HPC/Intro\\_to\\_IJC\\_HPC.html](https://ijcbit.github.io/Workshops/Trainings/HPC/Intro_to_IJC_HPC.html)
- Teams: ijcbioinfo

Questions?

